

Team Members:

Mengxing Chen (chenm32)

Anny Kong (yk57)

Yuqi Huang (yh73)

Xinrong Zhao (zhaox29)

Nullness_Lite: An unsound option of the Nullness Checker with fewer false positives

https://github.com/weifanjiang/Nullness_Lite

1. Abstract

In this paper, we are going to focus on the Nullness Checker part of the Checker Framework. We create another version of Nullness Checker, Nullness_Lite, as an option users can choose before using the complete Nullness Checker. Nullness_Lite is more user-friendly since it requires fewer annotations than the complete Nullness Checker, but it is unsound while the Nullness Checker is sound. We will focus on four features of the current Nullness Checker, and modify/disable each of them in order to reduce the number of false positives and the number of annotations needed.

We compare our Nullness_Lite with other checkers which can be used to find NullPointerException. Specifically, we choose a unit testing framework, JUnit4, which is widely used in industry and unannotated by any nullness bug detector or type system. We run each checker on JUnit4, and record the number of true positives, the number of false positives, the number of true positives the checker does not reveal, and the number of annotations needed to run this checker on JUnit4. We treat the complete Nullness Checker as the ground truth in order to record the number of true positives each checker does not reveal since the complete Nullness Checker is sound.

2. Problems & Motivation

Null is so significant to software development that almost all the programming languages allow programmers to write null. But null has been the cause of countless troubles in the history of software development. Tony Hoare, the inventor of null, calls it “the billion-dollar mistake” [10]. More than that, since null is the main reason for the infamous null pointer error, there is hardly a programmer who is not troubled by null. According to research, the null pointer exceptions (NPE) are reported as the most frequent bug in Java programs [31]. Null pointer dereferences are not only frequent [32], but also catastrophic [33]. Though Java already provides an infrastructure for exceptions, the current state of the language is just a partial solution. Professor John Sargeant from Manchester School of Computer Science also said “of the things which can go wrong at runtime in Java programs, null pointer exceptions are by far the most common” [3].

As developers seek a world where no NPEs are raised, some nullness bug detectors, such as FindBugs [11], NullAway [24], and the Nullness Checker of the Checker Framework [4], are emerging as a supplement of Java’s weak type system that does not support compile-time nullness checking. Because there is a correlation between popularity and star counts on GitHub [34], we noticed that there are about 2.1k stars for NullAway [24], 0.5k stars for FindBugs [23], and 361 stars for the Checker Framework [25]. We also noticed the Checker Framework and FindBugs are popular in industries. For instance, the Checker Framework is used on hundreds of projects at Google [43], and FindBugs is used on all projects

in Amazon. As we realized the Nullness Checker of the Checker Framework, the one with the sound analysis, is more popular in industries, and might be too complicated for less skilled developers on the large open-source platform to use, we had an idea: add an option with fewer false positives and make the sound Nullness Checker easier to use. As an initial step and evaluation of this idea, we built a new Nullness_Lite option on the Nullness Checker of the Checker Framework, the one that is sound but complicated to use. It is unsound but still effective in the static analysis of nullness errors. And it can provide a faster and easier option for Java developers who would like to get a compile-time nullness analysis on their source code, but hesitate to spend time running full verification tools like the Nullness Checker.

3. Recent Solutions & Related Works

NullAway is an annotation-based nullness bug detector which is mainly built to detect bugs in Android projects, its command-line based version also works on non-Android programs. It uses its own type system to detect nullness bugs [30]. It is fast: “the build-time overhead of running NullAway is usually less than 10%” [24]. Although it is built as a plugin of ErrorProne [17], it actually benefits most of the users since ErrorProne is widely used in industry. However, NullAway has its limitations. For example, it cannot check code using generics [16] and null assertions [18].

IDEs like IntelliJ [14] and Eclipse [15] also provide annotation-based null analysis. And they have their own libraries for annotations. IntelliJ supports two annotations, `@Nullable` and `@NotNull`. Eclipse support three annotations, `@NonNull`, `@Nullable`, and `@NonNullByDefault`. Also, IntelliJ has a functionality called “Infer Nullity” which automatically add `@Nullable` and `@NotNull` in the project. Also, they provide dataflow analysis which runs in the background so that when the user types in the program, they statically check for (possible) null-related errors.

FindBugs [27] is another nullness bug detector. It analyzes nullness bugs using heuristic based pattern matching which is a code idiom that has a high probability of being an error [41]. Also, FindBugs is powerful in that it can directly analyze the bytecode of a program [29], so the users do not even need the source code of their program in order to use FindBugs.

The Nullness Checker of Checker Framework is a sound, pluggable type checker [9], which, unlike those bug detectors introduced above, aims at detecting all nullness bugs. It requires users to add annotations into their code as machine-checked documentation that its type-based dataflow analysis, “a technique to statically derive information about the dynamic behavior of a program” [19], can take advantage of to give a precise error report. An experiment of checking the Lookup program [35] shows that the Nullness Checker has the stronger analysis than FindBugs by successfully detecting all 9 true positives of nullness bugs while FindBugs missed all [4].

Although the Nullness Checker introduces more false positives than other nullness bugs detectors that we mentioned above, customers still find it easy to use. The Nullness Checker has local type inference [26], so that if customers do not specify the annotations for a variable then it will automatically infer annotations for the variable and issue errors if no annotations allow the program to type-check. There are also some type inference tools customers can use as an extra step before checking their programs with the Nullness Checker. Since the unannotated types are analyzed by the Nullness Checker as `@NonNull` by default [26], customers can use the AnnotateNullable tool [36] of Daikon [37] to add `@Nullable`, although sometimes customers need to write additional `@Nullable`, it reduces the burden of customers to add annotations for using the Nullness Checker.

However, since “the Checker Framework is designed to value soundness over limiting false warnings” [4], some assumptions for nullness bugs analysis in the Nullness Checker are so conservative that

```
private volatile Test test;
// other methods in between
public JUnit38ClassRunner(Test test) {
    // [initialization.fields.uninitialized] test
    super();
    setTest(test);
}
// other methods in between
private void setTest(
    @UnderInitialization JUnit38ClassRunner this, Test test) {
    this.test = test;
}
```

(Figure 1. False positive example of Uninitialized Error)

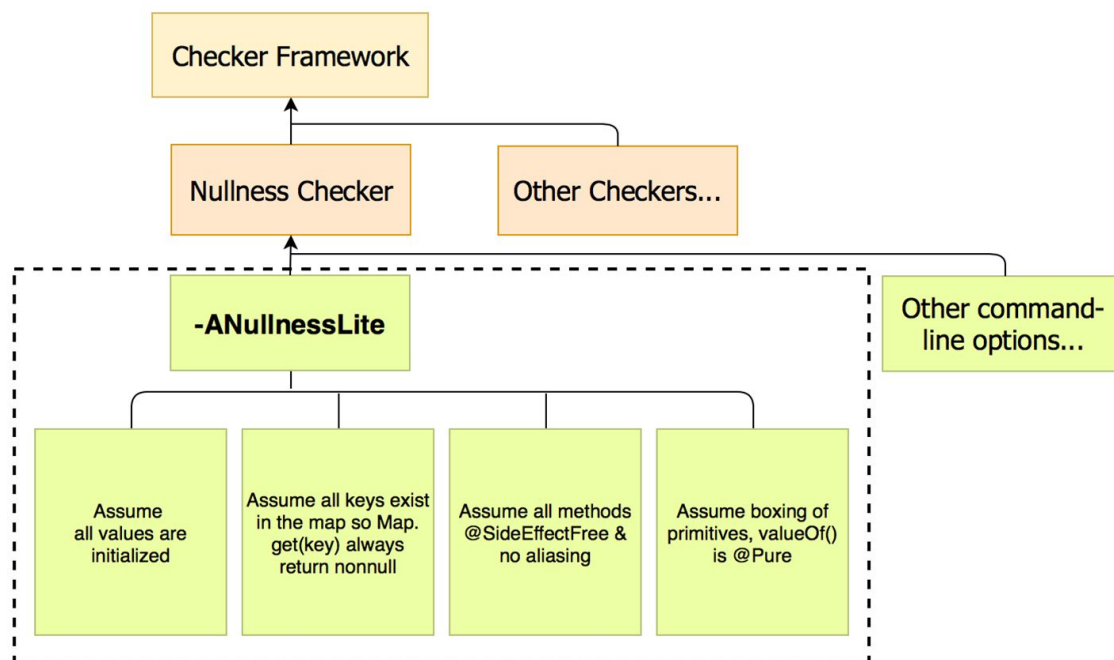
overwhelms the users by raising too many false positives. One false positive example raised is the JUnit38ClassRunner class [39] from JUnit4 [38]. Shown in figure 1, the Nullness Checker issues a field uninitialized error of the private instance variable test in JUnit38ClassRunner’s constructor, because it analyzes test to be @NonNull but test is not

literally assigned by a @NonNull value in the constructor. However, the constructor takes a @NonNull Test as the parameter and assign it to the test field in the helper method setTest. Therefore, test is actually initialized in the JUnit38ClassRunner class, a false positive. We will examine more about why the Nullness Checker raise this error and also other kinds of false positives in next section. After all, the false positives raised by the Nullness Checker require extra effort from users to verify the them manually.

4. Approaches

4.1 Hypothesis

Considering the suggestions given by Michael Ernst, one of the Checker Framework lead maintainers, and listed in the Nullness bug detector section of GSoC ideas 2018 [22], we attempted to build on the current Nullness Checker and add a new option ---- Nullness_Lite, as shown in figure 2.



(figure 2. Nullness_Lite Outline)

As an extension to the original Nullness Checker, `Nullness_Lite` is lite, but incomplete and unsound nullness checker. We keep the original sound checker, and provide a new possibility for users to start small and advance to full verification in the future, rather than having to start out doing full verification.

The `Nullness_Lite` option can be enabled via providing a command line argument `"-ANullnessLite"`. And it has four differences from the current Nullness Checker: disable the Initialization Checker and the Map Key Checker that the nullness analysis based on; modify some assumptions of the dataflow analysis and the behaviors of boxing primitives.

Our goal is to reduce the number of false positives. When testing with an open source program, Google Collections, the Nullness Checker was found to report 362 false positives while only reporting 9 true errors [9], which is obviously a significant trouble for bug fixing. The following four features were suggested as an implementation of this idea. After trying with each feature, we see the potential possibility for these features to produce a large number of false warnings, an annoying thing for programmers to deal with. They have to look into each warning and suppress each by a manual proof that it is not a true positive. In the following section, we give examples for some possible existing problems with false positives, which may have time costly effects, in the Nullness Checker's internal implementation:

- **Assume all values are initialized:**

For `@NonNull` values, the Nullness Checker assumes that all instance variables are initialized literally inside the constructors and all static class variables are initialized at their declarations. If violating either of the assumptions above, the Nullness Checker simply issue fields uninitialized errors no matter whether these fields actually raise NPEs or not.

The conservative assumptions are good to eliminate the bugs caused by uninitialized fields, but sometimes are impractical. Back to the `JUnit38ClassRunner` example in `JUnit4` shown in figure 1, we realized that developers sometimes decide to initialize `@NonNull` instance fields in helper methods of constructors in real world projects. Although the fields are actually initialized, the Nullness Checker still issues the false positive that can confuse the developers when debugging.

Therefore, we have the hypothesis that the unsound `Nullness_Lite` option can reduce the false positives by suppressing the uninitialized errors.

- **`Map.get` returns `@NonNull` result:**

The Map Key Checker which the Nullness Checker depends on tracks which values are keys for which maps. The Nullness Checker uses its `@KeyFor` annotation to determine if a value is a key for a given map — that is, to indicate whether `map.containsKey(value)` would evaluate to `true`. If variable `v` has type `@KeyFor("m") ...`, then the value of `v` is a key in Map `m`. That is, the expression `m.containsKey(v)` evaluates to `true` [27].

The `Map.get` method is annotated with `@Nullable` in the annotated jdk files, therefore assumed to be returning nullable values except some cases. The `@KeyFor` annotation is checked by a Map Key Checker that the Nullness Checker invokes. This annotation enables the Nullness Checker to treat calls to `Map.get` more precisely by refining its result to `@NonNull` if the following two conditions are satisfied [27]: 1) `mymap`'s values are all non-null; that is, `mymap` was declared as `Map<KeyType, @NonNull ValueType>`. Since `@NonNull` is the default type, it need not be written explicitly. 2) `mykey` is a key in `mymap`; that is,

`mymap.containsKey(mykey)` returns `true`. This fact could be expressed to the Nullness Checker by declaring `mykey` as `@KeyFor("mymap") KeyType mykey`.

However, the Map Key analysis is not strong enough to cover many cases when `Map.get` also returns a non null result. In particular, it is unable to analyze codes outside the scope of method where the variable `key` is declared. As `Map.get` method is assumed to be `@Nullable`, key-value pairs added elsewhere, say added in other methods, are not recognized by the Nullness Checker and will cause it to produce a false warning. For example, figure 3 is an example of such false positives in JUnit4. While we run the Nullness checker on the code, it gives a warning, while it should not since the key `validateWithAnnotation` is in the map `VALIDATORS_FOR_ANNOTATION_TYPES`. The variable `validateWithAnnotation` is added to the map `VALIDATORS_FOR_ANNOTATION_TYPES` when it is absent in method

```
public AnnotationValidator createAnnotationValidator(ValidateWith validateWithAnnotation) {  
    ...  
    AnnotationValidator annotationValidator = validateWithAnnotation.value().newInstance();  
    VALIDATORS_FOR_ANNOTATION_TYPES.putIfAbsent(validateWithAnnotation, annotationValidator);  
    return VALIDATORS_FOR_ANNOTATION_TYPES.get(validateWithAnnotation); // OK but a warning  
    ...  
}
```

(Figure 3. False positive example of the Map Key checker)

`putIfAbsent()` right before `get()` being called. Additionally, the corresponding value `annotationValidator` is not null which is ensured by `newInstance()`. Since the method `putIfAbsent()` is not within the scope of the function `createAnnotationValidator()`, Map Key Checker, the subchecker of the Nullness Checker, could not recognize its existence and assumes that it is not in the map `VALIDATORS_FOR_ANNOTATION_TYPES`, while it is actually in the map at the time when `get(validateWithAnnotation)` is called. As a result, it gives a false warning.

As demonstrated in the example above in practice, the Map Key Checker's analysis cannot always infer a true warning for `Map.get`, and there is a potential for more keys being added outside the current method than within the current method, therefore false warnings may reduce if we disable the Map Key Checker. We will make Nullness_Lite to assume for every call to `Map.get(key)`, the given key exists in the map and always returns a `@NonNull` result.

- **Assume all methods `@SideEffectFree` and no aliasing:**

The Nullness Checker considers two situations that potentially invalidate the dataflow facts. One is method invocation. The other is updating one field which potentially causes a change of its aliasing.

The Nullness Checker analyzes the former using the side effect analysis [26]. When a `@Nullable` variable is refined to `@NonNull`, the Nullness Checker will invalidate the `@NonNull` fact if any method having access to this variable is called. Yet the assumption is conservative because methods are considered "dangerous" even if they do not actually modify the fields. Thus, the Nullness Checker sometimes raise false positives for the real world projects because it is not uncommon for developers to write code that can cause this kind of false warnings.

One example is the `ComparisonCompactor` class [40] also from JUnit4, partial code showed in figure 4. The Nullness Checker issues an error at line 29, that the instance variable `fExpected`

is a `@Nullable` string but the method `compactString(String source)` requires the

```
21 @SuppressWarnings("deprecation")
22 public String compact(String message) {
23     if (fExpected == null || fActual == null || areStringsEqual()) {
24         return Assert.format(message, fExpected, fActual);
25     }
26
27     findCommonPrefix();
28     findCommonSuffix();
29     String expected = compactString(fExpected);
30     String actual = compactString(fActual);
31     return Assert.format(message, expected, actual);
32 }
```

(Figure 4. False positive example of Method Side Effect Error)

reassign `fExpected` to be null (code is not shown in figure 4). Thus, we draw the conclusion that `fExpected` is `@NonNull` at line 29. Yet the Nullness Checker invalidates the `@NonNull` fact after the method call at line 27 simply because it has the access to `fExpected`.

For the latter, when assigning null to some field of a variable, the Nullness Checker checks whether other variables under the same scope can be its aliases and invalidate the `@NonNull` fact of their fields as well. However, it conservatively assumes that any two objects can be aliases

```
import org.checkerframework.checker.nullness.qual.Nullable;

public class FWExample {
    public static class Node {
        public @Nullable Node next;
        public int val;

        public Node(int val) {
            this.val = val;
            this.next = null;
        }
    }

    public void foo(Node a) {
        Node b = new Node(0);
        if (a.next != null) {
            b.next = null;
            a.next.toString(); // False Positive Warning
        }
    }
}

zhaox29@ubuntu:~/403$ javac -processor nullness FWExample.java
FWExample.java:18: error: [dereference.of.nullable] dereference of
possibly-null reference a.next
    a.next.toString(); // False Positive Warning
    ^
1 error
```

(figure 5. A False Positive Example of Aliasing)

- **Boxing of primitives to be not `@Pure`:**

The Nullness Checker has most boxed classes' `valueOf(primitiveType)` method such as `Integer.valueOf(int)`, `Character.valueOf(char)`, `Short.valueOf(short)` being annotated as `@SideEffectFree` rather than `@Pure` in the annotated jdk for its soundness. The `valueOf` methods in wrapper classes always (`Boolean`, `Byte`) or sometimes (`Character`, `Integer`, `Short`) return an interned result as stated in JLS 5.1.7 because it does not guarantee that the boxed primitive is always the same object [28].

parameter to be `@NonNull`. However, this error is a false positive of invalidation of dataflow. First of all, `fExpected` is refined to be `@NonNull` after the if statement at line 23, where we exit the whole method if `fExpected` is null. Although having access to `fExpected`, the methods called at line 27 & 28 do not

of each other if one is the same type or the subtype of the other. In this way, although revealing all aliasing bugs, the Nullness Checker sometimes raises false positives for variables of the same type but not aliases of each other. We did not find the real world example for this kind of false warnings so far, but one “naive” example is shown in figure 5.

Due to reasons above, we predict that although making the Nullness Checker unsound, the `Nullness_Lite` option can reduce the false positives by both assuming all methods are `@SideEffectFree` and no aliasing allowed.

However, for type `int` between -128 and 127 inclusive, it will generate false warnings. JLS states “when the value `p` being boxed is an integer literal of type `int` between -128 and 127 inclusive, and let `a` and `b` be the results of any two boxing conversions of `p`, it is always the case that `a == b`” [28]. But since

```
void showIntegerValueOfFalseWarnings() {
    if (foo(Integer.valueOf(127)) != null) {
        foo(Integer.valueOf(127)).toString(); // OK but a warning
    }
}

@Pure
@Nullable String foo(int b) {
    return "nonnull";
}
```

(Figure 6. False positive example of the boxing of primitives)

the Nullness Checker assumes it to be only equals by `equals()` for being sound, it produces a false warning while passing an `int` between -128 and 127. As we have not encountered a real-world example so far, in the contrived example shown in figure 6,

`foo()` is expected to return the same string since it is annotated with `@Pure`. However, the Nullness Checker gives a false warning since it assumes two calls to `Integer.valueOf(127)` to return two different strings.

The false positives also happen when the boxing `Integer.valueOf(128)` is passed in as a parameter to some method, say `foo()`, then in this case, the method `foo()` will only return null if it distinguishes `Integer.valueOf(128)`. Hence, as long as the call of `foo` with the boxing of 128 returns nonnull, another call to `foo()` with the same boxing of 128 will also be nonnull if the method does not care about inputs or does not distinguish `Integer.valueOf(128)`. Then we will propose to assume that the JVM is always interning integers. Hence Nullness_Lite will assume the boxing of primitives always returns the same object on every call by replacing the original annotation `@SideEffectFree` with `@Pure`.

Although we are deliberately giving up soundness instead of purely improving the Nullness Checker, we provide users with an additional faster and easier-to-use option. With our unsound option, users can reach another point in the design space where fewer annotations are required and fewer false positives are produced. Our hypothesis states with the trade-off of soundness, the Nullness Checker with the Nullness_Lite option enabled should produce fewer false warnings and require fewer annotations, therefore faster and more usable for programmers.

4.2 Command-line User Interface

Nullness_Lite has a command line interface. To use Nullness_Lite checker, users will add a command line argument “-ANullnessLite” when using Nullness Checker. Note that “-ANullnessLite” option will only work for the Nullness Checker. The behavior is undefined if the option is passed to other checkers. All original commands in Checker Framework will still work. Figure 7 describes an example running Nullness_Lite through `javac`:

```
zhaox29@ubuntu:~/jsr308/checker-framework$ javac -processor nullness -ANullnessLite <MyFile.java>
```

(Figure 7. Nullness_Lite’s Command Line Usage)

5. Evaluation & Experiments

This section shows our evaluation work, the comparison of the nullness-related errors reported on a real world project by the Nullness_Lite option, the Nullness Checker, NullAway, FindBugs, IntelliJ, and Eclipse. We evaluated our project on the open-source Java Program, JUnit4 [38], which is a unit testing framework which is widely used in industry and unannotated by any of the nullness bug detector or type system. The code segment we analyzed excludes the test code in JUnit4, because it does not reflect the way developers would implement in real world projects.

The evaluation result table is shown in the figure 8. For reproducibility of the table, clients can follow the reproduction section in the user manual of our Nullness_Lite repository.

Checkers/Features	Bugs Revealed	Bugs Not Revealed	False Positives	Annotations Added
The Nullness_Lite Option	18	0	86	304
- All variables initialized	18	0	90	305
- <code>Map.get()</code> returns <code>@NonNull</code>	18	0	92	306
- No aliasing + all methods <code>@SideEffectsFree</code>	18	0	94	307
- <code>BoxedClass.valueOf()</code> are <code>@Pure</code>	18	0	95	307
The Nullness Checker	18	0	95	307
NullAway (using IntelliJ's Infer Nullity)	3	15	1	1285
NullAway (using annotations required for the Nullness Checker)	3	15	0	307
NullAway (using annotations required for the Nullness_Lite)	3	15	0	304
IntelliJ	0	18	1	0
IntelliJ (Infer Nullity)	3	15	77	15 ¹
FindBugs	0	18	0	0
Eclipse	0	18	3	0

(Figure 8. Evaluation Result on JUnit4)

5.1 Evaluation for the Nullness_Lite option and the Nullness Checker

We analyzed the Nullness_Lite option, each feature of the Nullness_Lite option and the Nullness Checker by the following steps:

1. Compile all source files with the checker we evaluate, with annotated JDK provided by the Checker Framework.
2. If some false positives can be reduced by simply adding annotations provided by the Nullness Checker, then leave brief comment why we added annotations here. Then add the annotations and

¹ Since we are evaluating JUnit4 based on the annotations added by IntelliJ's Infer Nullity, 15 is the number of annotations we changed rather than added. The number of annotations added by Infer Nullity is 1116.

repeat step 1 (because adding annotations sometimes results in new errors). If no false positives can be reduced by that, then proceed to step 3.

3. Analyzed the rest errors and leave a comment, which consists of the type of the error, either it is true or false positive, and a brief reasoning. An error is analyzed as an true positive (bug revealed) if the JUnit4 documentation lacks the information of nullness assumptions so that clients can write code to raise NPEs or the internal implementation raises NPEs. We provide examples for true positives related to clients' code. An error is analyzed as an false positive if the all situations to raise NPEs are impossible according to the JUnit4 documentation and implementation.

From the evaluation result shown in figure 8, the Nullness_Lite reduced 9 false positives related to the features being disabled compared to the Nullness Checker. Among these false positives, we found 5 reduced by "All variables initialized" feature, 3 reduced by "Map.get() returns @NonNull" feature, and 1 reduced by "all methods @SideEffectsFree" feature. Although we didn't find any errors related to features "No aliasing" and "BoxedClass.valueOf() are @Pure" in JUnit4, we reported it on our table as a reference for developers who want to extend our project in the future. We would also like to highlight that the JUnit4 is a special case where the Nullness_Lite option revealed all 18 true positives even though it is unsound.

5.2 Evaluation for IntelliJ

We evaluate the nullness checker of IntelliJ in 2 ways: 1. Run its nullness checker on JUnit4 which does not have any annotations; 2. Run its nullness checker on JUnit4 which has annotations added by Infer Nullity, a functionality that automatically examines the project and then adds @Nullable and @NotNull in the source code of the project.

We found only 1 false positive in the non-annotation version of JUnit4. And we cannot reduce this false positive by adding annotations in the source code. Detailed analysis is included in comments in the source code.

Infer Nullity introduced 1116 annotations into JUnit4. And then we ran IntelliJ's checker on this version of JUnit4. We examined each error found by IntelliJ's checker, and tried to eliminate some of them by changing some annotations added by Infer Nullity. It turns out that after changing 15 annotations, the number of errors are reduced to the largest extent: 77 errors are false positives and 3 are true positives. And we've included detailed analysis in comments in the source code.

It is noticeable that Infer Nullity adds this many annotations to JUnit4. Although the number of annotations added by Infer Nullity is much more than that added by either Nullness Checker or Nullness_Lite, Infer Nullity is effective in that it reduces around 15 false positives, compared to the Nullness Checker and Nullness_Lite. Also, as shown in figure 8, IntelliJ finds 3 true positives in the annotated version of JUnit4, which indicates that Infer Nullity is not powerful enough because Nullness Checker and Nullness_Lite can find more true positives with fewer annotations added, and true positives are what the users most care about.

5.3 Evaluation for NullAway

We used IntelliJ's infer nullity to add 1285 annotations automatically, consisting of @Nullable and @NotNull. NullAway reported 4 errors, where 3 of them were true positives. However, since we used annotations generated by IntelliJ, we might miss some annotations at current stage (since NullAway can check packages and files which are not fully annotated). In that case, since we already analyzed

annotations needed to run Nullness Checker and Nullness_Lite, we created another two branches to record the result of running NullAway with annotations required by Nullness Checker and Nullness_Lite, respectively. The Nullness Checker used 307 annotations, and NullAway found 3 errors (all of them are true positives). The Nullness_Lite used 304 annotations, and NullAway found the same 3 true positives.


5.4 Evaluation for FindBugs

We used findbugs-3.0.1 version. It does not need any change of its configuration since errors related to NPEs are placed in the same directory after fully analysing the project. Since FindBugs reported 8 null-related errors in JUnit4 test files and 0 errors in source files, we had the conclusion that FindBugs report 0 warnings and the results shown in the table above.

5.5 Evaluation for Eclipse

In order to run Eclipse's nullness analysis, we changed the compiler preference by checking "Null Pointer Access" and "Possible Null Pointer Access". Then, by rebuilding the project, Eclipse found 3 null-related errors. We tried to eliminate as many errors as possible by adding annotations into the source code, but it turns out that all the 3 errors found by Eclipse cannot be eliminated in this way. So the number of annotations added for Eclipse is zero. Also, we've classified that all 3 errors are false positives and then attached detailed analysis in the source code in order to justify our classification.

5.6 Evaluation Summary



(Figure 9. False positives v.s. True positives for nullness checkers)

Figure 9 shows the results according to the true and false positives. Our evaluation result does not imply that one checker is definitely better than others. Every checker has its own strengths and weaknesses,

which results in different design spaces. As a result, users should choose the tool that fit their situations best. For example, the Nullness Checker is a good choice when users value a good verification over the time consumed. The other bug detectors are good in the reversed situation. NullAway is good for users working for android projects; the IDEs like IntelliJ and Eclipse are good for users who want instant check while typing the code; FindBugs runs fast, has a low false positive rate, and can run on compiled bytecode; the Nullness_Lite option is in the middle ground of a good verification and fair time consumed.

6. Conclusion

This project includes implementing a lite version of the Nullness Checker of the Checker Framework and evaluating our lite version of the Nullness Checker along with several popular nullness checkers on an industrial Java project, JUnit4. Compared to Nullness Checker which we deem as the ground truth, our lite version of the Nullness Checker successfully reports fewer false positives but equal true positives on JUnit4. It provides evidence that the NullnessLite is effective as we expected and provides an upgrade path for users to full Nullness Checker analysis. Also, our evaluation is of value in that it includes almost all popular nullness checkers. By showing and comparing the data we have obtained from evaluating the checkers on JUnit4, we come to the conclusions of the relative advantages/disadvantages of these checkers compared to each other. Thus, our evaluation provides useful information for developers when they need to choose a proper checker in order to detect potential nullness bugs in their code.

7. Limitations and Future Work

There exists some limitations in our project. First, our experimental subject, JUnit4, is not the latest version of Java unit testing framework. It might not be well maintained. Since it is still a popular choice among developers based on star counts on GitHub, the information provided by our evaluation is still useful about the potential nullness bugs in JUnit4.

Also, the size of our evaluation object is limited and thus we wish we could have had more time to evaluate the checkers on more projects. With more projects, the result could be more generalizable and thus more convincing. Hence our future work will revolve around evaluating the checkers on a few more industrial projects. One avenue of future work is to do more evaluation on larger and well-maintained projects. This would provide more evidence on the effectiveness of the Nullness_Lite option. We could also consider more features to include in the NullnessLite option which allows faster and easier analysis for users.

8. Bibliography

1. "Interesting Facts about Null in Java." GeeksforGeeks, 14 Sept. 2017, www.geeksforgeeks.org/interesting-facts-about-null-in-java/.
2. Neumanns, Christian. "Why We Should Love 'Null'." *Why We Should Love 'Null' - CodeProject*, 18 Nov. 2014, www.codeproject.com/Articles/787668/Why-We-Should-Love-null.
3. Sargeant, John. "Null Pointer Exceptions." *Null Pointer Exceptions*, www.cs.man.ac.uk/~johns/npe.html.
4. Dietl, Werner, and Michael Ernst. "Preventing Errors Before They Happen The Checker Framework." *Preventing Errors Before They Happen The Checker Framework*, www.google.com/url?sa=t&rct=j&q=&esrc=s&source=web&cd=7&cad=rja&uact=8&ved=0ahUKEwiup6_V5bPaAhXKrFQKHW-gCxYQFghaMAY&url=https%3A%2F%2Fstatic.rainfocus.com%2Foracle%2Ffoo_w17%2Fsess%2F1492901668615001brln%2FFPF%2F2017-10-02%2520CF%2520%40%2520JavaOne_1507012791774001WJ2t.pdf&usg=AOvVaw3mAtzExTzYm6gr3sCn0cXb.

5. Osman, Haidar, et al. "Tracking Null Checks in Open-Source Java Systems." *2016 IEEE 23rd International Conference on Software Analysis, Evolution, and Reengineering (SANER)*, 2016, doi:10.1109/saner.2016.57.
6. Admin. "Java NullPointerException - Reasons for Exception and How to Fix?" *The Java Programmer*, 5 June 2017, www.thejavaprogrammer.com/java-nullpointerexception/.
7. Dobolyi, Kinga. "Changing Java's Semantics for Handling Null Pointer Exceptions."
8. Maciej Cielecki, Je,drzej Fulara, Krzysztof Jakubczyk, and Jancewicz. Propagation of JML non-null annotations in java programs. In *Principles and practice of programming in Java*, pages 135–140, 2006.
9. Dietl, Werner, et al. "Building and Using Pluggable Type-Checkers." *Proceeding of the 33rd International Conference on Software Engineering - ICSE '11*, 2011, doi:10.1145/1985793.1985889
10. Hoare, Tony. "Null References: The Billion Dollar Mistake." *InfoQ*, Tony Hoare, 25 Aug. 2009, www.infoq.com/presentations/Null-References-The-Billion-Dollar-Mistake-Tony-Hoare.
11. Ayewah, Nathaniel, et al. "Using Static Analysis to Find Bugs." *IEEE Software*, vol. 25, no. 5, 2008, pp. 22–29., doi:10.1109/ms.2008.130.
12. Sridharan, Manu. "Engineering NullAway, Uber's Open Source Tool for Detecting NullPointerExceptions on Android." *Uber Engineering Blog*, 15 Feb. 2018, eng.uber.com/nullaway/.
13. Hovemeyer, David, et al. "Evaluating and Tuning a Static Analysis to Find Null Pointer Bugs." *The 6th ACM SIGPLAN-SIGSOFT Workshop on Program Analysis for Software Tools and Engineering - PASTE '05*, 2005, doi:10.1145/1108792.1108798.
14. "Infer Nullity." *JetBrains*, www.jetbrains.com/help/idea/inferring-nullity.html.
15. *Help - Eclipse Platform*, help.eclipse.org/kepler/index.jsp?topic=/org.eclipse.jdt.doc.user/tasks/task-using_null_annotations.htm.
16. Sridharan, Manu. "Support Models of Generic Types · Issue #54." *GitHub, Uber/NullAway*, 6 Nov. 2017, github.com/uber/NullAway/issues/54.
17. Google. "Error Prone." *Error Prone*, 2017, errorprone.info/.
18. kevinzetterstrom. "Support for null assertions · Issue #122 · Uber/NullAway." *GitHub*, github.com/uber/NullAway/issues/122.
19. Abel, Andreas, et al. "Dataflow Framework for Checker Framework." *Dataflow Framework for Checker Framework*, courses.cs.washington.edu/courses/cse501/10au/JavaDFF.pdf.
20. Knizhnik, Konstantin, and Cyrille Artho. "About Jlint." *Jlint - Find Bugs in Java Programs*, jlint.sourceforge.net/.
21. Pmd. "PMD." *PMD*, 29 Apr. 2018, pmd.github.io/.
22. "GSoC Ideas 2018." *Checker Framework Organization: GSoC Ideas 2018*, rawgit.com/typetools/checker-framework/master/docs/developer/gsoc-ideas.html.
23. Findbugsproject. "Findbugsproject/Findbugs." *GitHub*, 23 Sept. 2017, github.com/findbugsproject/findbugs.
24. Uber. "Uber/NullAway." *GitHub*, github.com/uber/NullAway.
25. Typetools. "Typetools/Checker-Framework." *GitHub*, github.com/typetools/checker-framework.
26. *The Checker Framework Manual: Custom Pluggable Types for Java*, checkerframework.org/manual/#map-key-checker.
27. "FindBugs™ - Find Bugs in Java Programs." *FindBugs™ - Find Bugs in Java Programs*, findbugs.sourceforge.net/.
28. "Chapter 5. Conversions and Contexts." *Lesson: All About Sockets (The Java™ Tutorials > Custom Networking)*, docs.oracle.com/javase/specs/jls/se8/html/jls-5.html#jls-5.1.7.
29. "FindBugs™ Fact Sheet." *FindBugs™ - Find Bugs in Java Programs*, findbugs.sourceforge.net/factSheet.html.
30. "Engineering NullAway, Uber's Open Source Tool for Detecting NullPointerExceptions on Android." *Uber Engineering Blog*, 15 Feb. 2018, eng.uber.com/nullaway.
31. Maciej Cielecki, Je,drzej Fulara, Krzysztof Jakubczyk, and Jancewicz. Propagation of JML non-null annotations in java programs. In *Principles and practice of programming in Java*, pages 135–140, 2006.

32. Arthur van Hoff. The case for java as a programming language. *IEEE Internet Computing*, 1(1):51–56, 1997.
33. V Vipindegep and Pankaj Jalote. List of common bugs and programming practices to avoid them. Technical report, Indian Institute of Technology, Kanpur, March 2005.
34. Papamichail, Michail, et al. “User-Perceived Source Code Quality Estimation Based on Static Analysis Metrics.” 2016 IEEE International Conference on Software Quality, Reliability and Security (QRS), 2016, doi:10.1109/qrs.2016.22.
35. *Lookup*, 10 May 2018, types.cs.washington.edu/plume-lib/api/plume/Lookup.html.
36. The Daikon Invariant Detector User Manual, plse.cs.washington.edu/daikon/download/doc/daikon.html#AnnotateNullable.
37. Ernst, Michael D., et al. “The Daikon System for Dynamic Detection of Likely Invariants.” *Science of Computer Programming*, vol. 69, no. 1-3, 2007, pp. 35–45., doi:10.1016/j.scico.2007.01.015.
38. junit-team. “JUnit-Team/junit4.” *GitHub*, github.com/junit-team/junit4.
39. junit-team. “JUnit-Team/junit4.” *GitHub*, github.com/junit-team/junit4/blob/master/src/main/java/org/junit/internal/runners/JUnit38ClassRunner.java
40. junit-team. “JUnit-Team/junit4.” *GitHub*, github.com/junit-team/junit4/blob/master/src/main/java/junit/framework/ComparisonCompactor.java.
41. “Analysis Tool Report.” <http://www.cs.cmu.edu/~aldrich/courses/654/tools/simulacrum-findbugs-06.pdf>
42. “JUnit API.” *JUnit 4*, junit.org/junit4/javadoc/latest/.
43. “Checker Framework - 2017.” *Google*, Google, summerofcode.withgoogle.com/archive/2017/organizations/6349316859887616/.

Appendix

In the repository of the Nullness_Lite option, we provide users with the instructions to reproduce the results we got in figure 8 and our analysis details in branches of our forked JUnit 4 repository (<https://github.com/NullnessLiteGroup/junit4>) listed as follows.

The Checker for Evaluation	Branch Name
The Nullness_Lite Option	annos_nl_all
--All variables initialized	annos_nl_init
--No aliasing + all methods @SideEffectsFree	annos_nl_inva
--Map.get () returns @NonNull	annos_nl_mapk
--BoxedClass.valueOf () are @Pure	annos_nl_boxp
The Nullness Checker	annos_nc_all
IntelliJ	intellij1
IntelliJ with Infer Nullity	intellij2
Eclipse	eclipse
FindBugs	findbugs
NullAway with annotations added by Infer Nullity	Nullaway_Intellij
NullAway with annotations required for the Nullness Checker	Nullaway_nc
NullAway with annotations required for NullnessLite	Nullaway_nl

Total Hours Spent: 182